

Introduction to Parallel and Distributed Processing

Shared Memory

Jaroslav 'Jaric' Zola


<http://www.jzola.org/>

Suggested Reading

- Cilk+ reference guide
<https://software.intel.com/en-us/node/522579>
- General resources <https://www.cilkplus.org/>

Shared Memory Architecture

- A type of parallel architecture where multiple processors share main memory (e.g. modern multi-core processors)
- Software perspective: threads share address space



Why Shared Memory

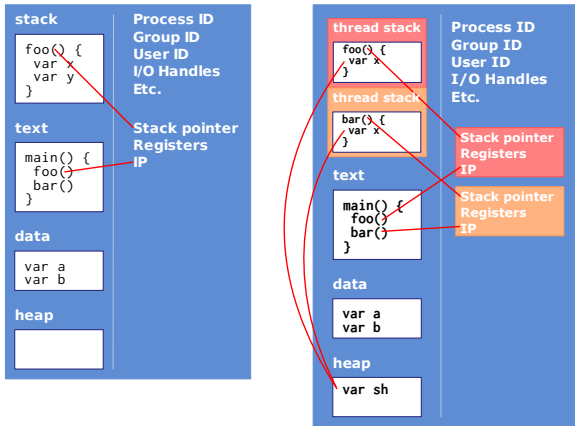
- Currently the most accessible type of parallel architecture
- Always at hand, why not use it
- Perceived easy parallelism

Why Shared Memory



Programming Shared Memory

- Create multiple threads that communicate by reading/writing shared variables




Programming Shared Memory

- Working with threads is too low-level for practical purposes
scheduling, synchronization
- At the high level two types of parallelism:
 - Task parallelism (MIMD): pool of threads executes tasks,
each task is a set of instructions on some data
 - Data parallelism: pool of threads acts on different parts
of data with the same set of instructions
- Ideally, we want to describe our tasks and data
transformations, without worrying about low-level realization


Fork-Join Pattern

- General strategy to express and execute task parallelism
- Main model for OpenMP, Cilk+, Intel TBB, etc.



Stencil Pattern

- Basic data parallel pattern: data structured (e.g. in 1D, 2D), update to data element depends on its neighborhood



Stencil

Cilk/Cilk+ API

- Provide a minimal set of C/C++ language extensions to enable shared memory parallelism (task and data)
- Focus on a parallel algorithms not implementation
- Efficient realization left to a compiler and run-time
- Comes with a powerful work-stealing scheduler and lock-free hyperobjects
- Deterministic Cilk+ code has serial semantics

Cilk+ Support

- Open source version in GCC 4.8 to 6.9, Clang/LLVM some support
- Commercial release by Intel (Parallel Studio XE, etc.)
- Using in GCC:

```
1| g++ -std=c++11 -fcilkplus -O3 ...
```

Spawn and Sync

- The Cilk+ standard defines three new keywords: `cilk_spawn`, `cilk_sync` and `cilk_for`
- These keywords are all what is needed to fork and join tasks

```
1  int fib(int n) {  
2      if (n < 2) return n;  
3  
4      // create child task  
5      int a = cilk_spawn fib(n - 1);  
6      int b = fib(n - 2);  
7  
8      // wait for all child tasks in function  
9      cilk_sync;  
10  
11     return a + b;  
12 } // implicit cilk_sync
```

Spawn and Sync

- Arguments passed to “spawned” function are evaluated before invocation
- Pass-by-references and pass-by-pointer arguments must be valid at least till `cilk_sync`
- Implicit `cilk_sync` follows after destructors
- It is a programmer job to avoid race conditions more to follow

Simple Parallel Postorder

```
1  struct node {
2      node* left;
3      node* right;
4  };
5
6  node* build_tree() { ... }
7  int release_tree(node* p) { ... }
8  void process(const node* p) { ... }
9
10 void visit(const node* p) {
11     if (p->left != 0) visit(p->left);
12     if (p->right != 0) visit(p->right);
13     process(p);
14 }
15
16 int main(int argc, char* argv[]) {
17     node* tree = build_tree();
18     visit(tree);
19     return release_tree(tree);
20 }
```

Simple Parallel Postorder

```
1  struct node {
2      node* left;
3      node* right;
4  };
5
6  node* build_tree() { ... }
7  int release_tree(node* p) { ... }
8  void process(const node* p) { ... }
9
10 void visit(const node* p) {
11     if (p->left != 0) cilk_spawn visit(p->left);
12     if (p->right != 0) visit(p->right);
13     cilk_sync;
14     process(p);
15 }
16
17 int main(int argc, char* argv[]) {
18     node* tree = build_tree();
19     visit(tree);
20     return release_tree(tree);
21 }
```

More Interesting Example

- $[f1, l1)$ and $[f2, l2)$ are sorted ranges

```
1  template <typename InIter, typename OutIter>
2  void xxxxx(InIter f1, InIter l1, InIter f2, InIter l2,
3           OutIter res) {
4      while ((f1 != l1) && (f2 != l2)) {
5          *(res++) = (*f2 < *f1) ? *(f2++) : *(f1++);
6      }
7
8      std::copy(f1, l1, res);
9      std::copy(f2, l2, res);
10
11 } // xxxxx
```


More Interesting Example

- $[f1, l1)$ and $[f2, l2)$ are sorted ranges

```
1  template <typename InIter, typename OutIter>
2  void merge(InIter f1, InIter l1, InIter f2, InIter l2,
3             OutIter res) {
4      while ((f1 != l1) && (f2 != l2)) {
5          *(res++) = (*f2 < *f1) ? *(f2++) : *(f1++);
6      }
7
8      std::copy(f1, l1, res);
9      std::copy(f2, l2, res);
10
11 } // merge
```

More Interesting Example

```
1  template <typename InIter, typename OutIter>
2  void merge_dc(InIter f1, InIter l1, InIter f2, InIter l2,
3              OutIter res) {
4      if ((l1 - f1) + (l2 - f2) < 5) {
5          return merge(f1, l1, f2, l2, res);
6      }
7
8      InIter mid1, mid2;
9
10     if ((l2 - f2) < (l1 - f1)) {
11         mid1 = f1 + ((l1 - f1) >> 1);
12         mid2 = std::lower_bound(f2, l2, *mid1);
13     }
14     else {
15         mid2 = f2 + ((l2 - f2) >> 1);
16         mid1 = std::lower_bound(f1, l1, *mid2);
17     }
18
19     merge_dc(f1, mid1, f2, mid2, res);
20     merge_dc(mid1, l1, mid2, l2, res + (mid1 - f1) + (mid2 - f2));
21 } // merge_dc
```

More Interesting Example

```
1  template <typename InIter, typename OutIter>
2  void merge_dc(InIter f1, InIter l1, InIter f2, InIter l2,
3               OutIter res) {
4      if ((l1 - f1) + (l2 - f2) < MERGE_LIMIT) {
5          return merge(f1, l1, f2, l2, res);
6      }
7
8      InIter mid1, mid2;
9
10     if ((l2 - f2) < (l1 - f1)) {
11         mid1 = f1 + ((l1 - f1) >> 1);
12         mid2 = std::lower_bound(f2, l2, *mid1);
13     }
14     else {
15         mid2 = f2 + ((l2 - f2) >> 1);
16         mid1 = std::lower_bound(f1, l1, *mid2);
17     }
18
19     cilk_spawn merge_dc(f1, mid1, f2, mid2, res);
20     merge_dc(mid1, l1, mid2, l2, res + (mid1 - f1) + (mid2 - f2));
21 } // merge_dc
```

More Interesting Example

```
1  template <typename InIter, typename OutIter>
2  void merge_sr(InIter f1, InIter l1, InIter f2, InIter l2,
3               OutIter res) {
4      while ((l1 - f1) + (l2 - f2) >= MERGE_LIMIT) {
5          InIter mid1, mid2;
6          if ((l2 - f2) < (l1 - f1)) {
7              mid1 = f1 + ((l1 - f1) >> 1);
8              mid2 = std::lower_bound(f2, l2, *mid1);
9          }
10         else {
11             mid2 = f2 + ((l2 - f2) >> 1);
12             mid1 = std::lower_bound(f1, l1, *mid2);
13         }
14         cilk_spawn merge_sr(f1, mid1, f2, mid2, res);
15         res += (mid1 - f1) + (mid2 - f2);
16         f1 = mid1;
17         f2 = mid2;
18     }
19     merge(f1, l1, f2, l2, res);
20 } // merge_sr
```

Efficiency Analysis

- If n is the total number of elements to merge:

$$T_1 = 2T_1\left(\frac{n}{2}\right) + O(\log(n)) = \Theta(n)$$

$$T_\infty = T_\infty\left(\frac{n}{2}\right) + O(\log(n)) = \Theta(\log^2(n))$$

$$S_\infty = O\left(\frac{n}{\log^2 n}\right)$$

For Fun

- Build parallel sorting routine starting with parallel merge