

Introduction to Parallel and Distributed Processing

Shared Memory Challenges

Jaroslav 'Jaric' Zola

<http://www.jzola.org/>

Shared Memory Challenges

- Parallelism is “easy” when only memory-read operations are involved
- In realistic scenarios tasks have to communicate and have to be synchronized
- Synchronization almost always penalizes efficiency
- Lack of proper synchronization may cause non-deterministic behavior
- Or lead to deadlock/livelock

Race Conditions

- Two tasks try to perform operation on the same memory at the same time, and one involves memory-write

```
1 | int add(int n, int x[]) {  
2 |     int S = 0;  
3 |     // race on S  
4 |     cilk_for (int i = 0; i < n; ++i) S = S + x[i];  
5 |     return S;  
6 | }
```

Race Conditions

- X and Y are shared between tasks and initially $X = Y = 0$,
 x and y are task local

T1	$X = 1$	$x = Y$	x and y both can be 0 or 1
T2	$Y = 1$	$y = X$	

- Modern architectures/compiler can reorder memory access
 – (almost) no consistency guarantees

Mutual Exclusion

- At the thread level race conditions can be addressed via mutual exclusion
- A thread locks access to the critical section using mutex

```

1  std::mutex g_mtx;
2
3  int add(int n, int x[]) {
4      int S = 0;
5      // DO NOT TRY THIS AT HOME!!!
6      cilk_for (int i = 0; i < n; ++i) {
7          g_mtx.lock();
8          S = S + x[i];
9          g_mtx.unlock();
10     }
11     return S;
12 } // add
    
```

Mutual Exclusion

- Lock/unlock are sequentially consistent and appear as atomic
- Critical section is executed only by one thread at a time
- Mutex can be unlocked only by the thread that locked it
- Hence, it is easy to get disaster:

M_X and M_Y are two mutexes

T1	$M_X.lock$	$M_Y.lock$	$X = 1$	$x = Y$	$M_Y.unlock$	$M_X.unlock$
T2	$M_Y.lock$	$M_X.lock$	$Y = 1$	$y = X$	$M_X.unlock$	$M_Y.unlock$

Deadlocks/Livelocks

- Deadlock: two or more threads block waiting for each other (e.g. on mutex)
- Livelock: two or more threads switch states but remain deadlocked
- Easy to get trapped if not careful: do not lock more than one mutex, always lock in the same order, avoid mutexes if possible

Cilk+ Approach

- Mutual exclusion is too low level and does not guarantee sequential consistency
- Instead, provide hyper-objects to enable lock-free access to shared variables
- General idea: each task sees a “local” version of a global object, final view is achieved via reduction
- User can define new reducers: reducer is always a monoid, i.e. a set with associative operator and identity value

Cilk+ Reducer Example

- Simple sum:

```
1  int add(int n, int x[]) {  
2      cilk::reducer<cilk::op_add<int>> S(0);  
3      cilk_for (int i = 0; i < n; ++i) *S += x[i];  
4      return S.get_value();  
5  } // add
```

More Interesting Example

```
1 | cilk::reducer_list_append<int> out;
2 |
3 | struct node {
4 |     int x;
5 |     node* left;
6 |     node* right;
7 | };
8 |
9 | void process(const node* p) {
10 |     out->push_back(p->x);
11 | }
12 |
13 | void visit(const node* p) {
14 |     if (p->left != 0) cilk_spawn visit(p->left);
15 |     if (p->right != 0) visit(p->right);
16 |     process(p);
17 | }
```

For Fun

- Implement prefix scan using Cilk+