

Introduction to Parallel and Distributed Processing

Introduction to OpenMP

Jaroslav 'Jaric' Zola

<http://www.jzola.org/>

Suggested Reading

- OpenMP reference guide <http://openmp.org/wp/>
- OpenMP tutorial
<https://computing.llnl.gov/tutorials/openMP/>

OpenMP Standard

- Programming standard for shared memory maintained by the industry consortium, support for C/C++ and Fortran
- Built around compiler directives, supporting run-time library and environment variables
- Focus on simplification of thread-based parallelism
- Based on fork-join paradigm, new standard directly addresses data parallelism

OpenMP Support

- All modern compilers provide strong support for OpenMP
- Using in GCC:

```
1 | g++ -std=c++11 -fopenmp -O3 ...
```

OpenMP General Idea

- Instrument code with “pragmas”, compiler behavior modifiers
- Mark regions that should be executed by a team of threads
- Mark level of parallelism, shared variables and synchronization

```
1 | #include <iostream>
2 |
3 | int main(int argc, char* argv[]) {
4 |   #pragma omp parallel
5 |   {
6 |     std::cout << "hello world" << std::endl;
7 |   }
8 |   return 0;
9 | } // main
```

OpenMP Syntax Basics

- Most of the OpenMP directives are of the form:
*#pragma omp construct [clause [,] clause]...] new-line
structured-block | for-loop*
- `#pragma omp parallel` construct forms a team of threads

```
1 | #include <iostream>
2 | #include <omp.h>
3 |
4 | int main(int argc, char* argv[]) {
5 |     #pragma omp parallel num_threads(2)
6 |     {
7 |         #pragma omp parallel num_threads(2)
8 |         {
9 |             if (omp_get_thread_num() == 0) {
10 |                 std::cout << omp_get_num_threads();
11 |             }
12 |         }
13 |     }
14 |     return 0;
15 | } // main
```

Basic Environment Variables

- **OMP_NUM_THREADS**
Default number of threads to use in parallel regions
- **OMP_THREAD_LIMIT**
Total number of threads allowed
- **OMP_NESTED**
Enable/disable nested parallel regions

Work Sharing Constructs

- Work sharing constructs:
for, sections, single, master, task
- Must be encountered by all threads in a team
- Work is distributed over all threads
- A work-sharing construct does not spawn any additional threads

Data Sharing Constructs

- `shared` – variable is shared between threads
- `private` – variable is private to each thread, i.e. each thread has a local copy, which is not initialized, and not passed outside of parallel region
- `firstprivate` – like `private`, but initialized
- `lastprivate` – like `private`, but the original copy updated after the construct

Data Sharing Rules

- The `for` loop iteration variable is **private**
- Automatic variables inside the `parallel` construct are **private**
- Variables with heap allocated storage are **shared**
- Static data members are **shared**
- Static variables declared in the `parallel` construct are **shared**
- Constants of type without mutable member are **shared**

OpenMP Synchronization

- Threads are synchronized implicitly at the beginning and at the end of parallel region and work sharing constructs (exception `master` construct and `nowait` clause)
- Threads can be explicitly synchronized (`barrier`, `critical`, `atomic`)
- Synchronization implies memory flushes

OpenMP Synchronization

- **critical** at any point of time only one thread can be executing region (critical section)
- **atomic** works only with expression statements, allows compiler to exploit atomic operations supported by CPU
- **barrier** specifies a point in the execution where all threads in a team wait for each other

Code Examples

- x shared by all threads executing loop

```
1  #include <iostream>
2  #include <vector>
3
4  int main(int argc, char* argv[]) {
5      int N = 1024 * 1024 * 1024;
6      std::vector<int> x(N);
7
8      #pragma omp parallel for shared(x)
9          for (int i = 0; i < N; ++i) x[i] = i;
10
11     return 0;
12 }
```

Code Examples

- What will happen if x is not firstprivate

```
1  #include <iostream>
2  #include <vector>
3
4  int main(int argc, char* argv[]) {
5      int N = 1024 * 1024;
6      std::vector<int> x(N, -1);
7
8      #pragma omp parallel
9      {
10         #pragma omp for firstprivate(x) nowait
11         for (int i = 0; i < N; ++i) x[i] = i;
12
13         #pragma omp critical
14         {
15             std::cout << x[0] << std::endl;
16         }
17     }
18     return 0;
19 } // main
```

Code Examples

```
1  #include <iostream>
2  #include <vector>
3
4  int main(int argc, char* argv[]) {
5      int N = 1024 * 1024;
6      std::vector<int> x(N, -1);
7
8      #pragma omp parallel shared(x)
9      {
10         #pragma omp for nowait
11         for (int i = 0; i < N; ++i) x[i] = i;
12
13         #pragma omp barrier
14         std::cout << "wow" << std::endl;
15     }
16     return 0;
17 } // main
```

Parallel For

- `schedule` – how loop should be distributed (static, dynamic, guided, auto, runtime)
- `reduction(operator:var)` – use `var` to perform reduction with operator

```
1 | #include <iostream>
2 | #include <vector>
3 |
4 | int main(int argc, char* argv[]) {
5 |     double S = 0.0;
6 |     int N = 1024 * 1024 * 1024
7 |     std::vector<double> x(N, 0.1);
8 |
9 |     #pragma omp parallel for shared(x) schedule(auto) \
10 |        reduction(+:S)
11 |        for (int i = 0; i < N; ++i) S += x[i];
12 |
13 |     return 0;
14 | } // main
```


Parallel Single & Master

- **single** – one thread executes the block, other wait
- **master** – only master executes the block, other proceed
- **nowait** – removes barrier from given block

```
1  #include <iostream>
2  #include <omp.h>
3  #include <unistd.h>
4
5  int main(int argc, char* argv[]) {
6      #pragma omp parallel
7      {
8          #pragma omp single nowait
9          {
10             sleep(1);
11             std::cout << "single: " << omp_get_thread_num() << std::endl;
12         }
13         #pragma omp master
14         {
15             std::cout << "master: " << omp_get_thread_num() << std::endl;
16         }
17     }
18     return 0;
19 } // main
```

Parallel Sections

- Each section executed by one thread
- If more sections than threads some threads execute multiple sections
- If more threads than sections idle threads wait (unless `nowait`)

```
1 | #include <iostream>
2 | #include <omp.h>
3 |
4 | int main(int argc, char* argv[]) {
5 |     #pragma omp parallel sections
6 |     {
7 |         #pragma omp section
8 |         std::cout << "A: " << omp_get_thread_num() << std::endl;
9 |
10 |        #pragma omp section
11 |        std::cout << "B: " << omp_get_thread_num() << std::endl;
12 |    }
13 |    return 0;
14 | } // main
```

Fibonacci Again

```
1  #include <iostream>
2
3  int fib(int n) {
4      int i, j;
5      if (n < 2) return n;
6      #pragma omp parallel sections shared(i, j) num_threads(2)
7          {
8          #pragma omp section
9              i = fib(n - 1);
10         #pragma omp section
11             j = fib(n - 2);
12         }
13         return i + j;
14     } // fib
15
16     int main(int argc, char* argv[]) {
17         int x = fib(64);
18         return 0;
19     } // main
```

Parallel Tasks

- task construct defines a task to be executed by a thread
- `taskwait` waits for completion of all **child** tasks

```
1 | #include <iostream>
2 |
3 | int fib(int n) {
4 |     int i, j;
5 |     if (n < 2) return n;
6 |     #pragma omp task shared(i)
7 |     i = fib(n - 1);
8 |
9 |     #pragma omp task shared(j)
10 |    j = fib(n - 2);
11 |
12 |    #pragma omp taskwait
13 |    return i + j;
14 | } // fib
15 |
16 | int main(int argc, char* argv[]) {
17 |     #pragma omp parallel
18 |     #pragma omp single nowait
19 |     std::cout << fib(36);
20 |     return 0;
21 | } // main
```

Parallel Tasks

- `depend` provides mechanism for defining task dependencies
- `untied` makes task independent of threads

```
1  int fib(int n) {
2      int i, j, S;
3      if (n < 2) return n;
4
5      #pragma omp task shared(i) depend(out:i)
6          i = fib(n - 1);
7
8      #pragma omp task shared(j) depend(out:j)
9          j = fib(n - 2);
10
11     #pragma omp task shared(i, j, S) depend(in:i,j) untied
12         S = i + j;
13
14     #pragma omp taskwait
15     return S;
16 } // fib
17
18 int main(int argc, char* argv[]) {
19     #pragma omp parallel
20     #pragma omp single nowait
21         fib(36);
22     return 0;
23 } // main
```

For Fun

- Implement Cilk+ examples using OpenMP tasks