

# Introduction to Parallel and Distributed Processing

## Introduction to Apache Spark

Jaroslav 'Jaric' Zola

<http://www.jzola.org/>

## Suggested Reading

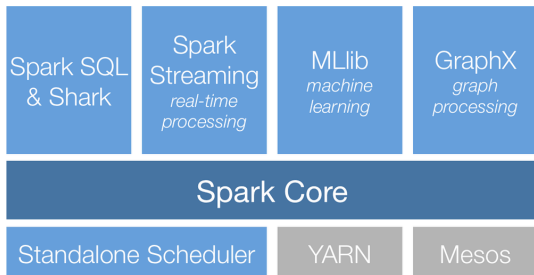
- Apache Spark <http://spark.apache.org/>
- “Learning Spark” Karu, H. et al. O'Reilly 2015

# Spark Overview

- Distributed memory computing platform extending Map/Reduce model
- Fairly unified stack working with different cluster management systems and storage systems
- Support for iterative processing and “shared variables”

# Spark Stack

- Core implementation in Scala, bindings to Java and Python
- Support for SQL, BSP model on graphs and real-time apps built on top



# General Idea

- Use a master-worker model, with “driver program” and executors
- Driver creates the execution context and keeps track of the data
- Executors do the actual work and store the data

# Spark Context

- Abstraction representing compute cluster
- Can be used to “configure” cluster (e.g. set environment)

```
1 | #!/usr/bin/python
2 | # ex01.py
3 |
4 | from pyspark import SparkContext
5 | import time
6 |
7 | if __name__ == "__main__":
8 |     sc = SparkContext(appName="Spark App")
9 |     time.sleep(120)
10 |    sc.stop()

1 | spark-submit --num-executors=8 --executor-cores 4 ex01.py
```

# Resilient Distributed Datasets

- Read-only, distributed collection of objects
- Programmer has control over partitioning and persistence
- RDD can be created either by transforming other RDD or from stable storage
- RDDs use lazy evaluation
- Fault tolerance is achieved via “lineage”

# Creating RDD

```
1 | #!/usr/bin/python
2 | # ex02.py
3 |
4 | from pyspark import SparkConf, SparkContext
5 | from random import random
6 |
7 | if __name__ == "__main__":
8 |     sc = SparkContext()
9 |     # not parallel
10 |     rn = sc.parallelize([random() for i in xrange(100000)])
11 |     sc.stop()
```



# Creating RDD

```
1  #!/usr/bin/python
2  # ex03.py
3
4  from pyspark import SparkConf, SparkContext
5  import sys
6
7  if __name__ == "__main__":
8      conf = SparkConf().setAppName("RDDcreate")
9      sc = SparkContext(conf = conf)
10     lines = sc.textFile(sys.argv[1])
11     sc.stop()
```

# Creating RDD

```
1 | #!/usr/bin/python
2 | # ex04.py
3 |
4 | from pyspark import SparkConf, SparkContext
5 | import json
6 | import sys
7 |
8 | if __name__ == "__main__":
9 |     sc = SparkContext()
10 |     rdd = sc.textFile(sys.argv[1]).map(lambda x: json.loads(x))
11 |     sc.stop()
```

# Transformations

- Create new RDD by applying some processing to already existing RDD
- Examples: flatMap, reduceByKey, filter, repartition

```

1 | #!/usr/bin/python
2 | # ex05.py
3 |
4 | from pyspark import SparkConf, SparkContext
5 | import sys
6 |
7 | def vertex(line):
8 |     l = line.split(" ")
9 |     v = [ int(x) for x in l ]
10 |    return [(v[0], 1), (v[1], 1)]
11 |
12 | if __name__ == "__main__":
13 |     conf = SparkConf().setAppName("RDDcreate")
14 |     sc = SparkContext(conf = conf)
15 |     lines = sc.textFile(sys.argv[1])
16 |     V = lines.flatMap(vertex) \
17 |         .reduceByKey(lambda a, b: a + b)
18 |     sc.stop()

```

# Actions

- Apply operation on RDD and return result to the driver, or write to stable storage
- Force evaluation of transformations
- Examples: `count`, `reduce`, `take`, `saveAsTextFile`

# Actions

```
1  #!/usr/bin/python
2  # ex06.py
3
4  from pyspark import SparkConf, SparkContext
5  import sys
6
7  def vertex(line):
8      l = line.split(" ")
9      v = [ int(x) for x in l ]
10     return [(v[0], 1), (v[1], 1)]
11
12  if __name__ == "__main__":
13     conf = SparkConf().setAppName("RDDcreate")
14     sc = SparkContext(conf = conf)
15     lines = sc.textFile(sys.argv[1])
16     V = lines.flatMap(vertex) \
17         .reduceByKey(lambda a, b: a + b)
18     n = V.count()
19     print n
20     V.saveAsTextFile(sys.argv[2])
21     sc.stop()
```

# Transformations and Actions

- In terms of implementation, action triggers DAG generation for RDD
- DAG reflects transformations
- Narrow transformations do not involve shuffling
- Wide transformations involve shuffling
- Narrow/wide transformations form stages
- Execution plan is created and passed to scheduler

# Persistence

- RDD is materialized when action is performed
- By default RDD is not preserved in the memory
- Unless `persist` method is invoked
- Type of storage and compression can be configured

```
1 | #!/usr/bin/python
2 | # ex07.py
3 |
4 | from pyspark import SparkConf, SparkContext
5 | import sys
6 |
7 | def edge(line):
8 |     v = [ int(x) for x in line.split(" ") ]
9 |     return (min(v[0], v[1]), max(v[0], v[1]))
10 |
11 | if __name__ == "__main__":
12 |     sc = SparkContext()
13 |     E = sc.textFile(sys.argv[1]).map(edge)
14 |     E.persist()
15 |     m = E.count()
16 |     print m
17 |     sc.stop()
```

# Passing Functions

- Functions/lambda's passed to transformations/actions and their closures are sent to executors whenever needed
- Be careful when using transformation/actions that modify local variables

```
1 | #!/usr/bin/python
2 | # ex08.py
3 |
4 | from pyspark import SparkConf, SparkContext
5 | import sys
6 |
7 | if __name__ == "__main__":
8 |     sc = SparkContext()
9 |     c = 0
10 |
11 |     # BAD IDEA!
12 |     def edge(line):
13 |         global c
14 |         c += 1
15 |         v = [ int(x) for x in line.split(" ") ]
16 |         return (min(v[0], v[1]), max(v[0], v[1]))
17 |
18 |     E = sc.textFile(sys.argv[1]).map(edge)
19 |     n = E.count()
20 |     sc.stop()
```



# Shared Variables

- Shared variables provide limited mechanism to accumulate/broadcast dependent data
- Accumulators do not trigger RDD evaluation!

```
1 | #!/usr/bin/python
2 | # ex09.py
3 |
4 | from pyspark import SparkConf, SparkContext
5 | import sys
6 |
7 | if __name__ == "__main__":
8 |     sc = SparkContext()
9 |     cacc = sc.accumulator(0)
10 |
11 |     def edge(line):
12 |         global cacc
13 |         cacc += 1
14 |         v = [ int(x) for x in line.split(" ") ]
15 |         return (min(v[0], v[1]), max(v[0], v[1]))
16 |
17 |     E = sc.textFile(sys.argv[1]).map(edge)
18 |     n = E.count()
19 |     print n, cacc
20 |     sc.stop()
```

# Shared Variables

```
1  #!/usr/bin/python
2  # ex10.py
3
4  from pyspark import SparkConf, SparkContext
5  import random
6  import sys
7
8  if __name__ == "__main__":
9      sc = SparkContext()
10     L = [random.random() for i in xrange(100000)]
11
12     # NOT EFFICIENT
13     def label(line):
14         global L
15         v = [int(x) for x in line.split(" ")]
16         return ((min(v[0], v[1]), max(v[0], v[1])), random.choice(L))
17
18     E = sc.textFile(sys.argv[1]).map(label)
19     E.saveAsTextFile(sys.argv[2])
20     sc.stop()
```

# Shared Variables

- Broadcasting

```
1 | #!/usr/bin/python
2 | # ex11.py
3 |
4 | from pyspark import SparkConf, SparkContext
5 | import random
6 | import sys
7 |
8 | if __name__ == "__main__":
9 |     sc = SparkContext()
10 |    L = sc.broadcast([random.random() for i in xrange(100000)])
11 |
12 |    def label(line):
13 |        global L
14 |        v = [ int(x) for x in line.split(" ") ]
15 |        return ((min(v[0], v[1]), max(v[0], v[1])), random.choice(L.value))
16 |
17 |    E = sc.textFile(sys.argv[1]).map(label)
18 |    E.saveAsTextFile(sys.argv[2])
19 |    sc.stop()
```

# For Fun

- Consider the problem of connected components over the list of edges. Implement the initialization step in Spark.