Introduction to Parallel and Distributed Processing

CUDA Reduce and Prefix

Jaroslaw 'Jaric' Zola
http://www.jzola.org/

Classic Problem

- Given vector $X = (x_1, \ldots, x_n)$ compute $\sum_{i=1}^{n} x_i$
- Smells like a simple 1D problem...

Direct Approach

```
1  __global__ void reduce0(int* gin, int* gout) {
2    extern __shared__ int sdata[];
3
4    int tid = threadIdx.x;
5    int i = blockIdx.x * blockDim.x + tid;
6
7    sdata[tid] = gin[i];
8    __syncthreads();
9
10   for (int s = 1; s < blockDim.x; s *= 2) {
11     if (tid % (2 * s) == 0) sdata[tid] += sdata[tid + s];
12     __syncthreads();
13   }
14
15   if (tid == 0) gout[blockIdx.x] = sdata[0];
16 } // reduce0
```

Direct Approach – Problem

Divergent branching:

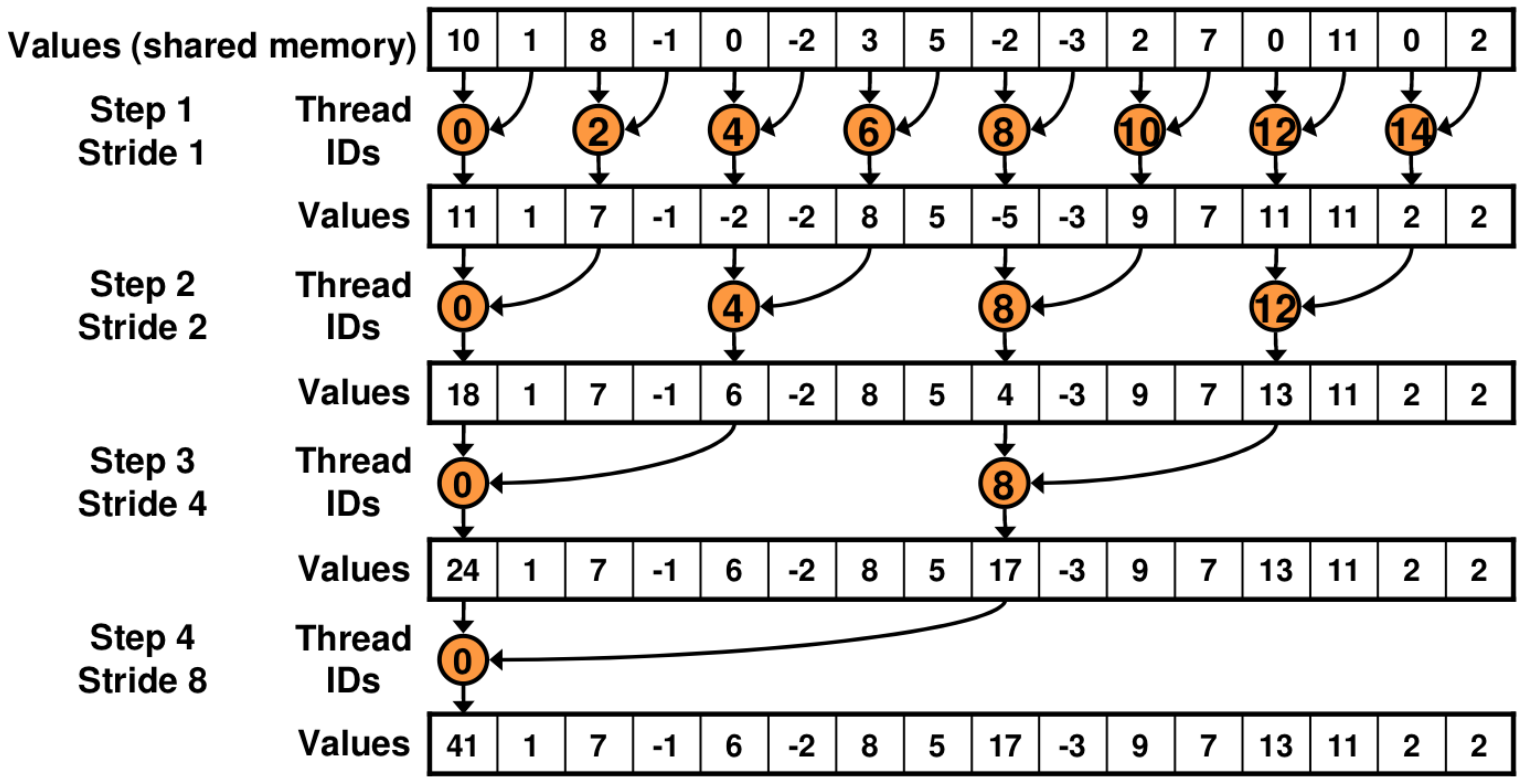

Better Approach

Sequential addressing:

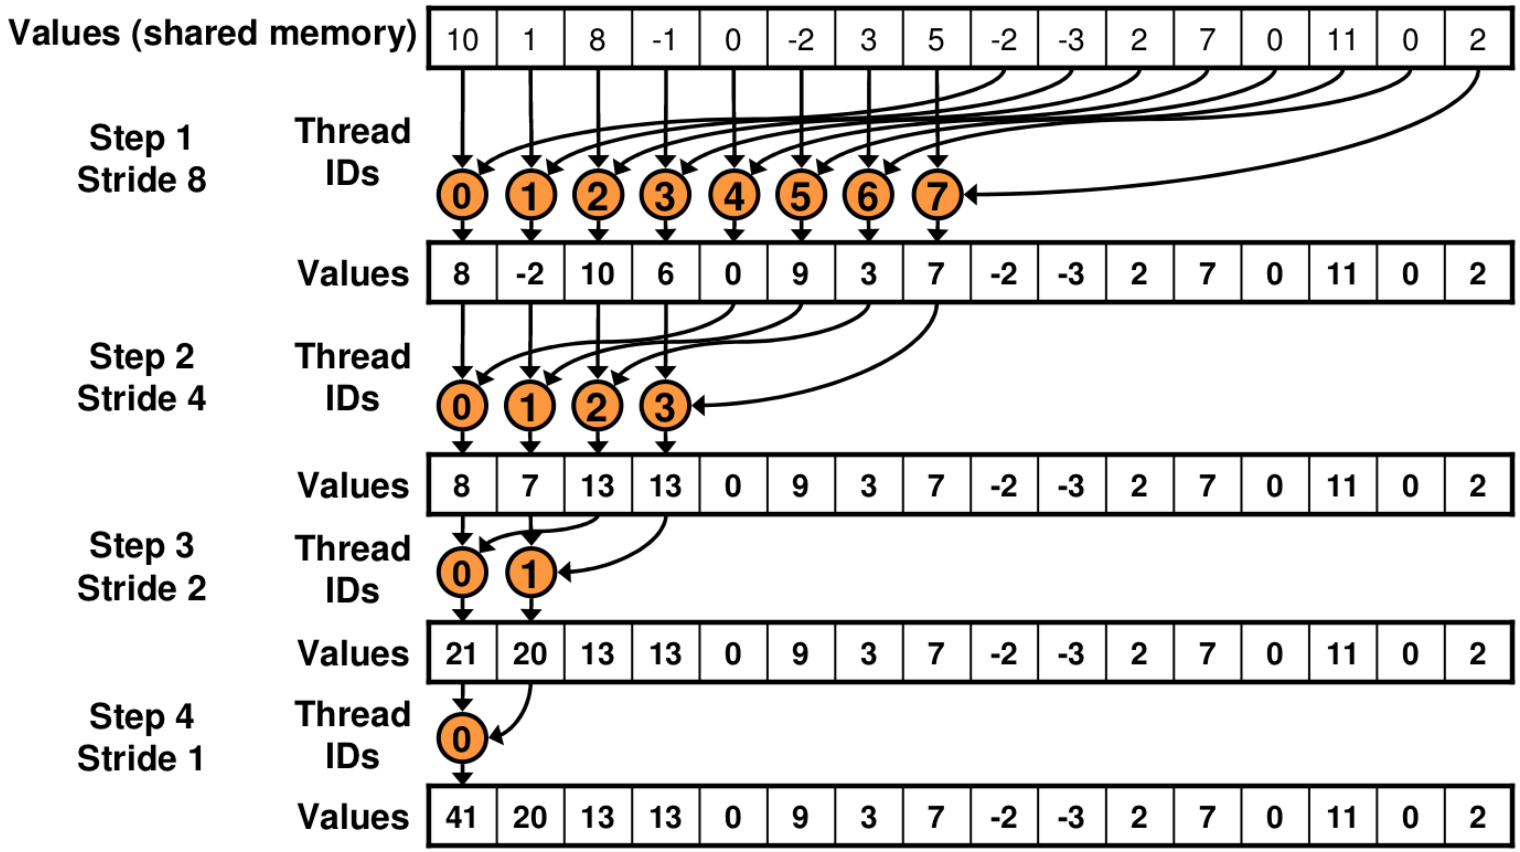

Sequential Addressing

```
1  __global__ void reduce1(int* gin, int* gout) {
2    extern __shared__ int sdata[];
3
4    int tid = threadIdx.x;
5    int i = blockIdx.x * blockDim.x + tid;
6
7    sdata[tid] = gin[i];
8    __syncthreads();
9
10   for (int s = blockDim.x >> 1; s > 0; s >>= 1) {
11     if (tid < s) sdata[tid] += sdata[tid + s];
12     __syncthreads();
13   }
14
15   if (tid == 0) gout[blockIdx.x] = sdata[0];
16 } // reduce1
```

Even Better Sequential Addressing

```
1  // we halve the number of blocks
2  __global__ void reduce2(int* gin, int* gout) {
3    extern __shared__ int sdata[];
4
5    int tid = threadIdx.x;
6    int i = blockIdx.x * (blockDim.x * 2) + threadIdx.x;
7
8    sdata[tid] = gin[i] + gin[i + blockDim.x];
9    __syncthreads();
10
11   for (int s = blockDim.x >> 1; s > 0; s >>= 1) {
12     if (tid < s) sdata[tid] += sdata[tid + s];
13     __syncthreads();
14   }
15
16   if (tid == 0) gout[blockIdx.x] = sdata[0];
17 } // reduce2
```

What About Parallel Prefix?

- Again, 1D problem PREFIX_SUM

Input: $[x_0, x_1, \ldots, x_{n-1}]$

1: **for** $d = 1 \ldots \log(n)$ **do**
2: **for** $j = 1 \ldots n$ **pardo**
3: **if** $j \geq 2^d$ **then**
4: $x[j] = x[j - 2^{d-1}] + x[j]$

What About Parallel Prefix?

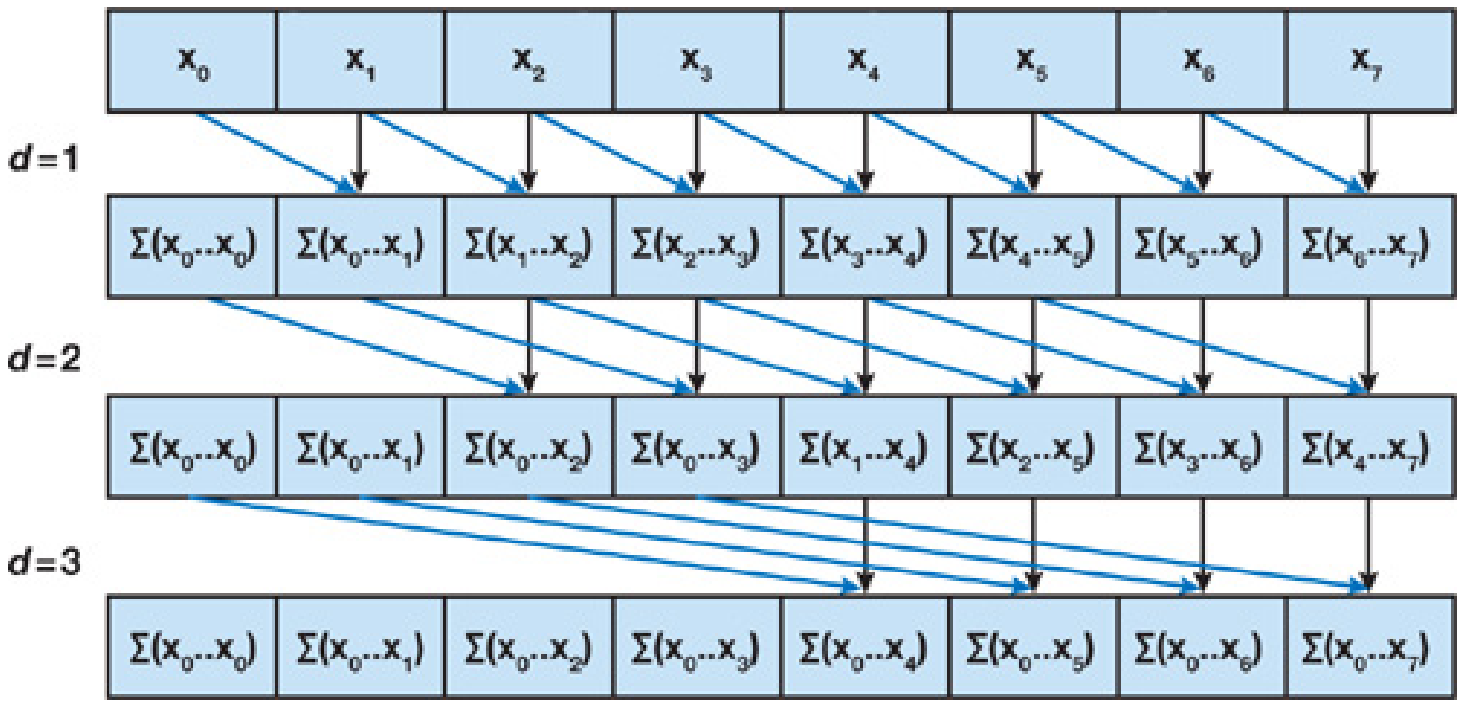

What About Parallel Prefix?

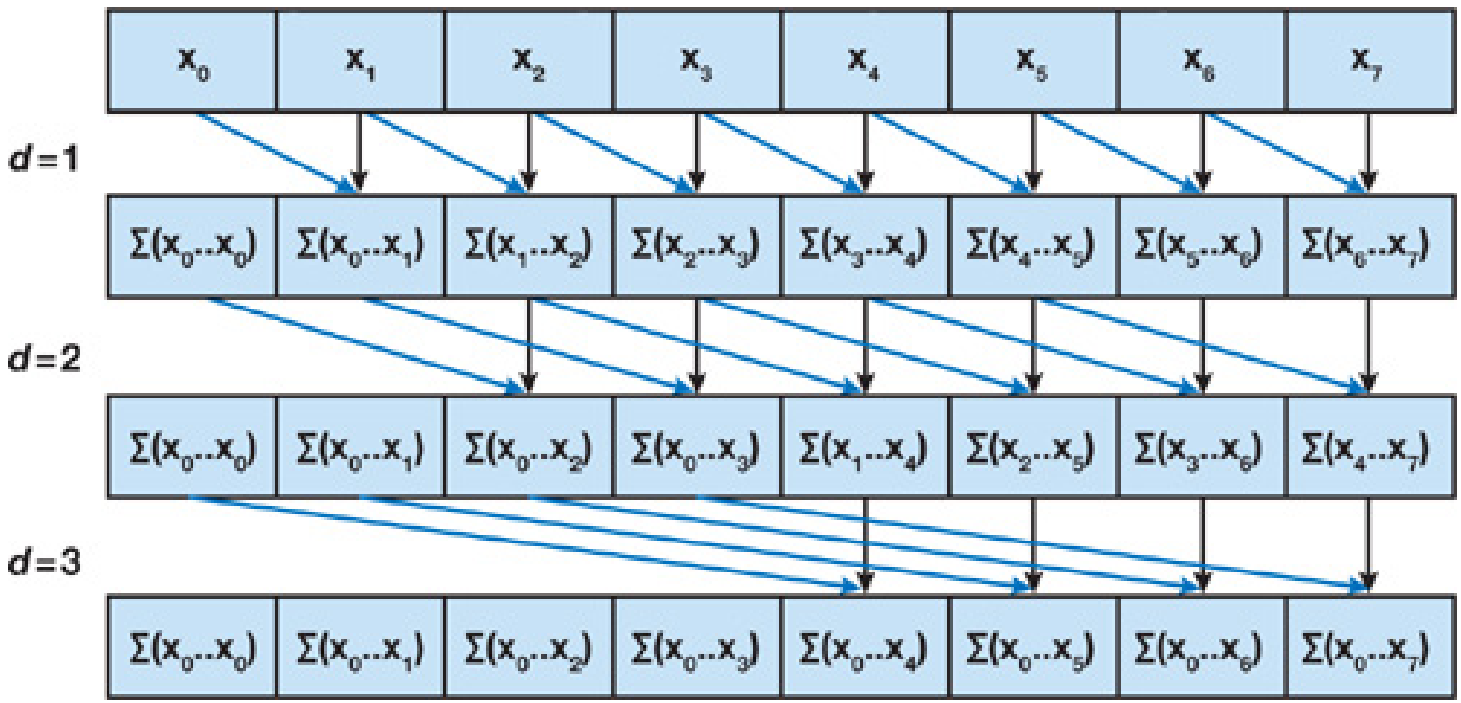

Problem: total work $O(n \log(n))$

Going Work Efficient

- We resolve to our old friend :-)
- Decompose problem into two phases: up-sweep (reduction) and down-sweep (update)

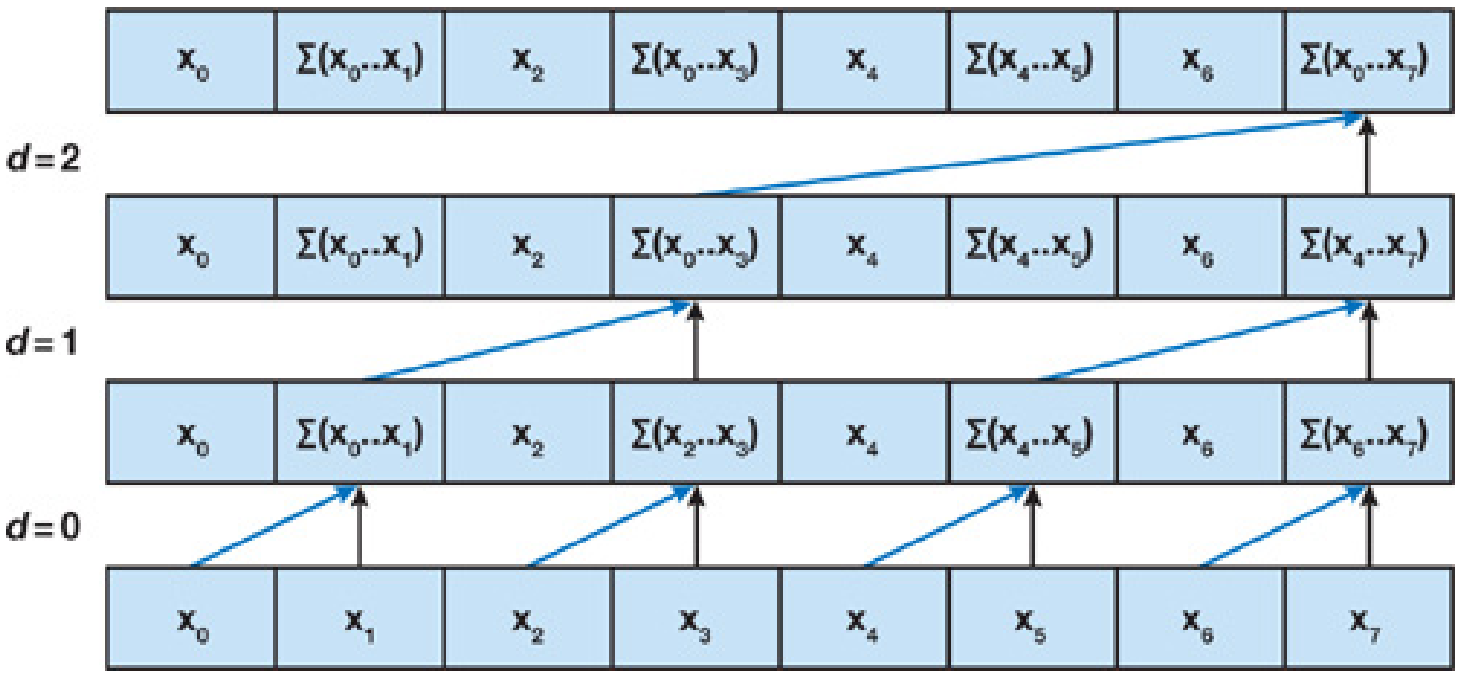

Going Work Efficient

- We resolve to our old friend :-)
- Decompose problem into two phases: up-sweep (reduction) and down-sweep (update)

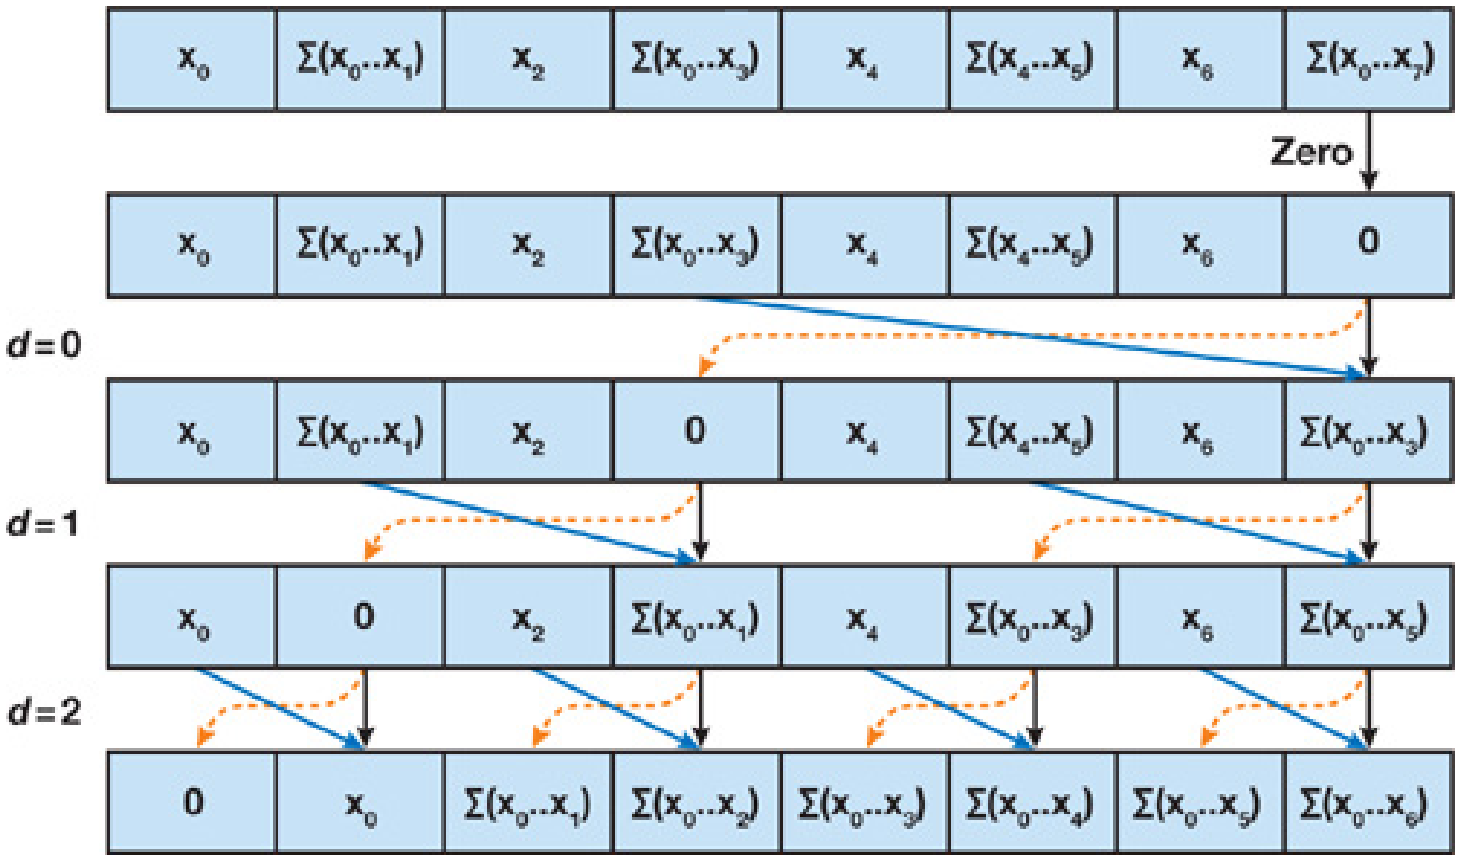

Working on Large Arrays

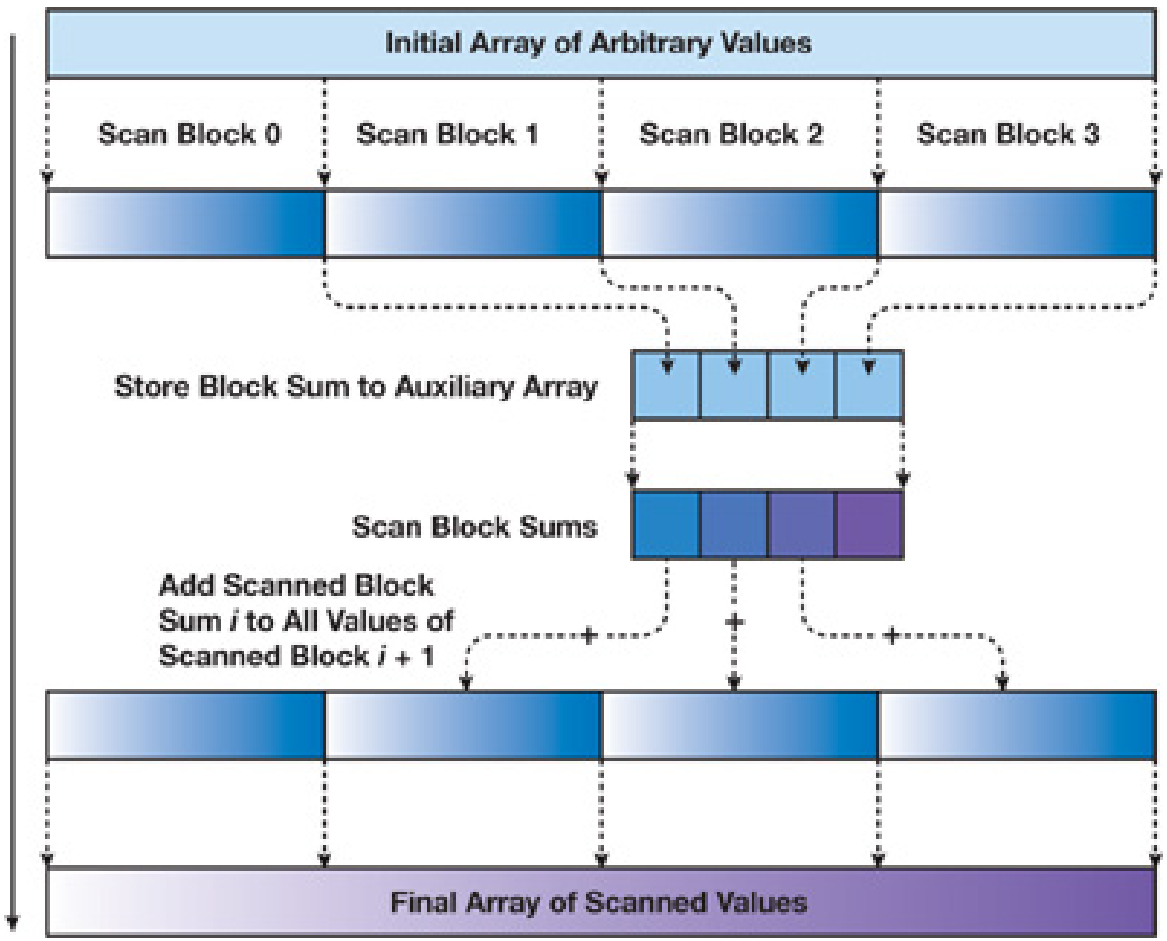

For Fun

- Modify the prefix code to work with large arrays.